

Beginning Data Exploration and Analysis with Apache Spark

GETTING STARTED WITH SPARK'S RESILIENT DISTRIBUTED DATASETS



Swetha Kolalapudi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Understanding the role of Spark in data analysis

Understanding RDDs and their characteristics

Installing Spark Standalone in a local environment

Loading data from a file

Reading data from an RDD



Big Data Borat

@BigDataBorat



Follow

In Data Science, 80% of time spent prepare data, 20% of time spent complain about need for prepare data.



There is no shortage of data
today, public or private

Public and Private

Comments

Reviews

Tweets

Blogs

**Government
Reports**

Purchases



Pageviews

Clicks



Messages

**Surveys/
Questionnaires**

Untold Opportunities



Draw insights



Apply Machine Learning

Doing Something Meaningful with Data

- Unstructured
- Messy
- Semantically complex



- Structured
- Clean
- Easy to consume

Data Processing

Data Processing Tasks

Parsing fields from text

Accounting for missing values

Identifying and investigating anomalies

Summarizing using tables and charts

Data Complexity

Messiness



Speed of Scaling

Small Data
Somewhat
Messy

Spreadsheets



Low data collection frequency


10s - 1000s of rows per day

Sometimes involves manual data collection

Many many files

Data Complexity

Messiness ↑



→
Speed of Scaling

Medium Data

**High Data
Integrity**

Databases



High frequency of collection


~100K rows per day

Programmatically collected

ACID properties

Data Complexity

Messiness ↑




→
Speed of Scaling

Big Data

**Very
messy**

Distributed
Computing



Very high frequency of data collection


Millions/Billions of rows per day

Files stored across a cluster of machines

Many many files (webpages, log files)

Data Complexity

Messiness ↑



→
Speed of Scaling

Tools for Data Munging



Excel



SQL
Python
Java



Hadoop
MapReduce
Spark

Spark

An engine for data processing and analysis



General Purpose



Interactive



**Distributed
Computing**

General Purpose



Exploring

Cleaning and Preparing

Applying Machine Learning

Building Data Applications

Spark

An engine for data processing and analysis



General Purpose



Interactive



**Distributed
Computing**

Interactive



```
len([1, 2, 5])  
>>> 3
```

REPL

Interactive



REPL

Read-Evaluate-Print-Loop

Interactive environments

Fast feedback

Spark

An engine for data processing and analysis



General Purpose



Interactive



**Distributed
Computing**

Distributed Computing



**Process data across a
cluster of machines**

Integrate with Hadoop

Read data from HDFS

Spark APIs

Scala

Python

Java

Almost all data is processed using

Resilient Distributed Datasets

Resilient Distributed Datasets

**RDDs are the main programming
abstraction in Spark**

Resilient Distributed Datasets

RDDs are in-memory collections of objects

**In-memory,
yet resilient!**

Resilient Distributed Datasets

**With RDDs, you can interact and
play with billions of rows of data**

**...without caring about
any of the complexities**

Spark is made up of a
few different components

Spark Core

The basic functionality of Spark
RDDs

Spark Core

Spark Core is just a computing engine

It needs two additional
components

Spark Core

A **Storage System**
that stores the
data to be
processed

A **Cluster Manager** to
help Spark run tasks
across a cluster of
machines

Spark Core

Storage
System

Cluster
Manager

**Both of these are plug
and play components**

Local file
system
HDFS

Storage
System

Spark Core

Storage
System

Cluster
Manager

**Built-in Cluster
Manager**

YARN

**Cluster
Manager**

Spark Core

Storage
System

Cluster
Manager

**Plug and Play makes it easy
to integrate with Hadoop**

A Hadoop Cluster

HDFS

For storage


YARN

**For managing
the cluster**

MapReduce

**For
computing**

A Hadoop Cluster



Use Spark as an
alternate/additional
compute engine

Installing Spark

Prerequisites

Java 7 or above

Scala

IPython (Anaconda)

Installing Spark

Download Spark binaries

Update environment variables

Configure iPython Notebook for Spark

Spark Environment Variables

SPARK_HOME

Point to the folder where Spark has been extracted

PATH

`$SPARK_HOME/bin` **Linux/Mac OS X**

`%SPARK_HOME%/bin` **Windows**

Configuring IPython

PYSPARK_DRIVER_PYTHON

ipython


PYSPARK_DRIVER_PYTHON_OPTS

“notebook”

Launch PySpark

> pyspark

```
Welcome to
```



```
version 1.6.1
```


```
Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)
```

```
SparkContext available as sc, HiveContext available as sqlContext.
```

```
>>> 
```

Launch PySpark


Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)


SparkContext available as sc, HiveContext available as sqlContext.

>>> 

This is just like a Python shell

Launch PySpark

```
Welcome to
```



```
version 1.6.1
```

```
Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)
```


```
SparkContext available as sc, HiveContext available as sqlContext.
```

```
>>> █
```

Use Python functions, dicts, lists etc

Launch PySpark


Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)


SparkContext available as sc, HiveContext available as sqlContext.

>>> 

**You can import and use any
installed Python modules**

Launch PySpark


Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)


SparkContext available as sc, HiveContext available as sqlContext.

>>> 

Launches by default in a local non-distributed mode

SparkContext

Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)


SparkContext available as sc, HiveContext available as sqlContext.

>>>

**When the shell is launched it
initializes a SparkContext**

SparkContext

Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)


SparkContext available as sc, HiveContext available as sqlContext.

>>>

The SparkContext represents a connection to the Spark Cluster

SparkContext

Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)


SparkContext available as sc, HiveContext available as sqlContext.

>>>

Used to load data into memory
from a specified source

SparkContext

Welcome to



version 1.6.1

Using Python version 2.7.11 (default, Jan 22 2016 08:29:18)

SparkContext available as sc, HiveContext available as sqlContext.

>>>

The data gets loaded into an RDD

Resilient
Distributed
Datasets

Partitions
Read-only
Lineage

**RDDs
represent data
in-memory**

Partitions

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi
3	Navdeep	25	Mumbai
4	Janani	35	New Delhi
5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

**Data is divided
into partitions**

**Distributed to
multiple
machines**

Partitions

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi
3	Navdeep	25	Mumbai
4	Janani	35	New Delhi
5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

**Data is divided
into partitions**

Partitions

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi
3	Navdeep	25	Mumbai
4	Janani	35	New Delhi
5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

**Data is divided
into partitions**

Partitions

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi

3	Navdeep	25	Mumbai
4	Janani	35	New Delhi

5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

**Distributed to
multiple
machines,
called nodes**

Partitions


1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi

3	Navdeep	25	Mumbai
4	Janani	35	New Delhi

5	Navdeep	25	Mumbai
6	Janani	35	New Delhi


**Distributed to
multiple
machines,
called nodes**

Partitions



**Nodes
process data
in parallel**

Partitions



Resilient
Distributed
Datasets

Partitions
Read-only
Lineage

Read-only

RDDs are immutable

Only Two Types of Operations



Transformation

**Transform into
another RDD**

Action

Request a result

Transformation

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi
3	Navdeep	25	Mumbai
4	Janani	35	New Delhi
5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

**A data set
loaded into
an RDD**

Transformation

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi
3	Navdeep	25	Mumbai
4	Janani	35	New Delhi
5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

**The user may
define a chain of
transformations
on the dataset**

Transformation

1	Swetha	30	Bangalore
2	Vitthal	35	New Delhi
3	Navdeep	25	Mumbai
4	Janani	35	New Delhi
5	Navdeep	25	Mumbai
6	Janani	35	New Delhi

1. Load data
2. Pick only the 3rd column
3. Sort the values

Transformation

- 1. Load data**
- 2. Pick only the 3rd column**
- 3. Sort the values**



Transformation

**Wait until a result is
requested before
executing any of
these transformations**



Only Two Types of Operations



Transformation

**Transform into
another RDD**

Action

Request a result

Action

**Request a
result using
an action**


- 1. The first 10 rows**
- 2. A count**
- 3. A sum**

Action

Data is processed only when the user requests a result


The chain of transformations defined earlier is executed

Lazy Evaluation



**Spark keeps a record
of the series of
transformations
requested by the user**

Lazy Evaluation



It groups the transformations in an efficient way when an Action is requested

Resilient
Distributed
Datasets

Partitions
Read-only
Lineage

Lineage

**When created, an RDD
just holds metadata**

1. A transformation

2. It's parent RDD

Lineage

**Every RDD
knows where it
came from**

RDD 2

Transformation 1

RDD 1

Lineage

Lineage can be
traced back all the
way to the source

RDD 2

Transformation 1

RDD 1

Lineage

Lineage can be traced back all the way to the source

RDD 2

Transformation 1

RDD 1

Load data

data.csv

Lineage

RDD 2

**When an action is
requested on an
RDD**

**All its parent RDDs
are materialized**

Transformation 1

RDD 1

Load data

data.csv

Lineage

Resilience

In-built fault tolerance

If something goes wrong,
reconstruct from source

Lazy Evaluation

**Materialize only when
necessary**

Summary

Understood the role of Spark in data analysis

Loading data from a file is a Spark transformation

Reading data from an RDD is a Spark Action

RDDs are partitioned, read-only collections which know their lineage

Spark transformations are lazily evaluated